



# Python as glue

Codes available on [GitHub](https://github.com/Paul-St-Young/thw-python-as-glue): <https://github.com/Paul-St-Young/thw-python-as-glue>

Yubo “Paul” Yang, THW-IL, 2018/04/17

Images from <https://pixabay.com/vectors>



# Step 1: Basic Interface

Choices:

f2py (numpy)

swig ([KMCLib](#))

ctype ([pyscf](#))

cython ([yt](#))

numba

boost.python

pybind11 (Eigen)

...

Criteria:

**easy** to write and debug

**compatible** with np.array

**separable** into native languages

Python can import  
*shared library*

f2py

```
NAME=example
all:
    f2py -c -m $(NAME) $(NAME).f90
clean:
    rm $(NAME).so
```

```
1 integer function add(i, j)
2   integer, intent(in) :: i, j
3   add=i+j
4 end function
```

pybind11

```
NAME=example
CXX=g++
OFLAGS=-O3
CFLAGS=$(OFLAGS) -shared -fPIC
OBJS=example.o

%.o: %.cpp
    $(CXX) $(CFLAGS) -c $<

all: $(OBJS)
    $(CXX) $(CFLAGS) $(OBJS) -o $(NAME).so

clean:
    rm *.o $(NAME).so
```

```
1 #include <pybind11/pybind11.h>
2
3 int add(int i, int j)
4 {
5     return i+j;
6 }
7
8 PYBIND11_MODULE(example, m)
9 {
10     m.def("add", &add, "add two integers");
11 }
12
```

<https://carbon.now.sh>

<https://team411.github.io/src2img>

## Step 2: Same Instruction Multiple Data (SIMD)

Python driver

```
1 import numpy as np
2 from forlib import example as fex
3 from cpplib import example as cex
4
5 def add(i, j):
6     return i+j
7
8 def check_correct(vec1, vec2):
9     pyvec = add(vec1, vec2)
10    cvec = cex.add(vec1, vec2)
11    fvec = fex.add(vec1, vec2)
12    print('cpp=python', np.allclose(cvec, pyvec))
13    print('fortran=python', np.allclose(fvec, pyvec))
14
```

('cpp=python', True)

('fortran=python', True)

('python/fortran', 1.689771011187475)

('python/cpp', 0.26767694027701155)

('python/python', 1.0473552503390706)

Timing: C++ is slowest

Fortran backend

```
1 subroutine add(vec1, vec2, vecout, n)
2   integer, intent(in) :: n
3   real*8, intent(in) :: vec1(n), vec2(n)
4   real*8, intent(out) :: vecout(n)
5   vecout = vec1+vec2
6 end subroutine
```

C++ backend

```
1 #include <pybind11/pybind11.h>
2 #include <pybind11/numpy.h>
3
4 namespace py=pybind11;
5
6 double add(double i, double j)
7 {
8     return i+j;
9 }
10
11 PYBIND11_MODULE(example, m)
12 {
13     m.def("add", py::vectorize(add), "add two integers");
14 }
```

# Step 3: Linear Algebra

Timing: numpy is fastest!

('python/fortran', 0.05105879678631556)

('python/cpp', 0.049230687039213586)

(numpy/numpy', 0.9886795048143053)

```
1 def matmul(mat, vec):  
2     return np.dot(mat, vec)
```

Fortran backend

```
1 subroutine mmul(mat, vec, vecout, m, n)  
2     integer, intent(in) :: m, n  
3     real*8, intent(in) :: mat(m, n), vec(n)  
4     real*8, intent(out) :: vecout(n)  
5     vecout = matmul(mat, vec)  
6 end subroutine
```

C++ backend

```
1 #include <pybind11/pybind11.h>  
2 #include <pybind11/eigen.h>  
3 #include <Eigen/Dense>  
4  
5 namespace py=pybind11;  
6 typedef Eigen::MatrixXd Matrix;  
7 typedef Eigen::VectorXd Vector;  
8  
9 Vector mmul(  
10     Eigen::Ref<const Matrix>& mat,  
11     Eigen::Ref<const Vector>& vec)  
12 {  
13     return mat*vec;  
14 }  
15  
16 PYBIND11_MODULE(example, m)  
17 {  
18     m.def("mmul", &mmul, "multiply matrix vector");  
19 }
```

## Step 4: General Computing (e.g. distance table)

$$r_{ij} = |\mathbf{r}_i - \mathbf{r}_j|$$

```
1 def disp_in_box(pos, lbox):
2     nint = np.round(pos/lbox)
3     return pos-nint*lbox
4
5 def get_dtable(pos, lbox):
6     natom = len(pos)
7     dtable = np.zeros([natom, natom])
8     for i in range(natom):
9         for j in range(natom):
10             dtable[i, j] = np.linalg.norm(
11                 disp_in_box(pos[i]-pos[j], lbox)
12             )
13     return dtable
```

```
1 subroutine distance_table(pos, lbox, natom, ndim, dtable)
2     integer, intent(in) :: natom, ndim
3     real*8, intent(in) :: pos(natom, ndim), lbox
4     real*8, intent(out) :: dtable(natom, natom)
5     integer i, j
6     real*8 drij(ndim)
7     do i=1,natom
8         do j=i+1,natom
9             drij = pos(i, :) - pos(j, :)
10            drij = drij - lbox*nint(drij/lbox)
11            dtable(i, j) = norm2(drij)
12        end do
13    end do
14 end subroutine
```

Timing: C++ and Fortran are both O(1000)  
times faster than Python for loops

('python/cpp', 734.6365928831605)

('python/fortran', 747.8374215630347)

('python/python', 1.0)

```
1 Eigen::Ref<Matrix> distance_table(
2     Eigen::Ref<const Matrix>& pos, double lbox)
3 {
4     const int natom = pos.rows();
5     const int ndim = pos.cols();
6     Matrix dtable = Matrix::Zero(natom, natom);
7     Vector drij(ndim);
8     for (int i=0; i<natom; i++)
9     {
10         for (int j=i; j<natom; j++)
11         {
12             drij = pos.row(i) - pos.row(j);
13             for (int idim=0; idim<ndim; idim++)
14             {
15                 drij(idim) -= lbox*std::round(drij(idim)/lbox);
16             }
17             dtable(i, j) = drij.norm();
18         }
19     }
20     return dtable;
21 }
```

## Step 4: General Computing (e.g. distance table)

$$r_{ij} = |\mathbf{r}_i - \mathbf{r}_j|$$

```
1 def disp_in_box(pos, lbox):
2     nint = np.round(pos/lbox)
3     return pos-nint*lbox
4
5 def get_dtable_vec(pos, lbox):
6     dptable = disp_in_box(
7     pos[:, np.newaxis] - pos[np.newaxis, :],
8     lbox)
9     dtable = np.linalg.norm(dptable, axis=-1)
10    return dtable
```

```
1 subroutine distance_table(pos, lbox, natom, ndim, dtable)
2     integer, intent(in) :: natom, ndim
3     real*8, intent(in) :: pos(natom, ndim), lbox
4     real*8, intent(out) :: dtable(natom, natom)
5     integer i, j
6     real*8 drij(ndim)
7     do i=1, natom
8         do j=i+1, natom
9             drij = pos(i, :) - pos(j, :)
10            drij = drij - lbox*nint(drij/lbox)
11            dtable(i, j) = norm2(drij)
12        end do
13    end do
14 end subroutine
```

Timing: vectorization saves some face  
but imposes other limitations

('pyvec/cpp', 6.621653701352949)

('pyvec/fortran', 6.754378146746862)

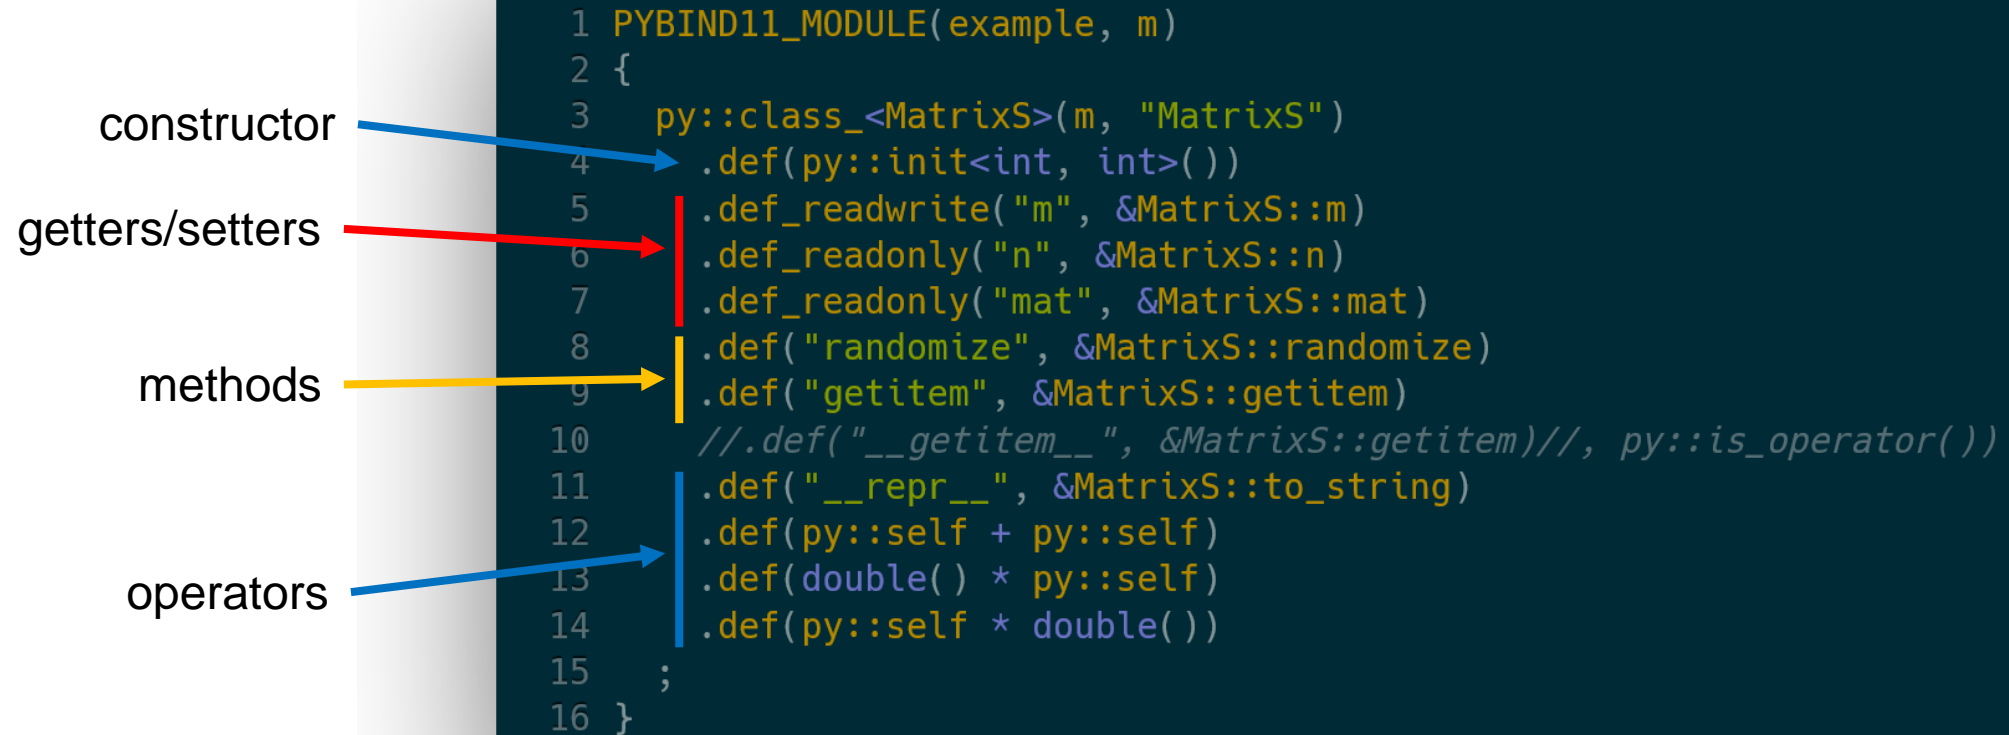
('pyvec/pyvec', 1.0)

```
1 Eigen::Ref<Matrix> distance_table(
2     Eigen::Ref<const Matrix>& pos, double lbox)
3 {
4     const int natom = pos.rows();
5     const int ndim = pos.cols();
6     Matrix dtable = Matrix::Zero(natom, natom);
7     Vector drij(ndim);
8     for (int i=0; i<natom; i++)
9     {
10         for (int j=i; j<natom; j++)
11         {
12             drij = pos.row(i) - pos.row(j);
13             for (int idim=0; idim<ndim; idim++)
14             {
15                 drij(idim) -= lbox*std::round(drij(idim)/lbox);
16             }
17             dtable(i, j) = drij.norm();
18         }
19     }
20     return dtable;
21 }
```

## Step 5: Classes

Binding to C++ class is straight-forward. Interface *borrowed* from boost.python.

Clean support for operator overloading (including pickling!)



The diagram illustrates the mapping between Python class features and their C++ counterparts in a Pybind11 module. On the left, labels with arrows point to specific lines of code in a C++ snippet. The labels are: 'constructor' (blue arrow to line 3), 'getters/setters' (red arrow to line 5), 'methods' (yellow arrow to line 9), and 'operators' (blue arrow to line 12). The code snippet is as follows:

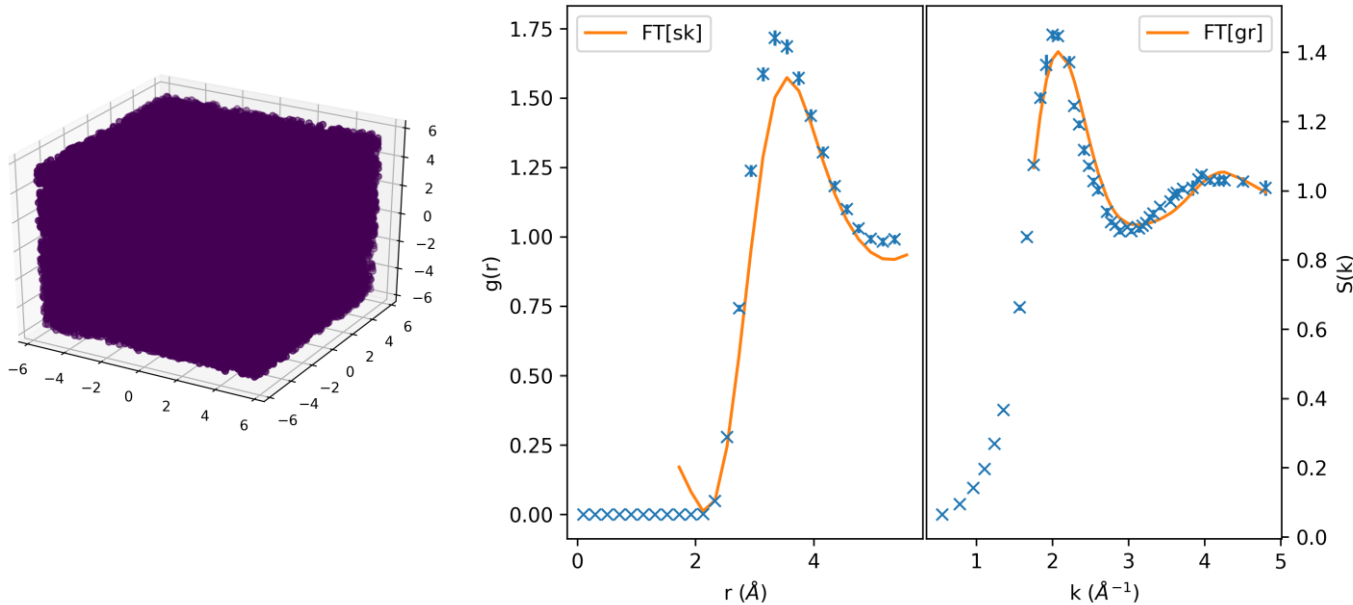
```
1 PYBIND11_MODULE(example, m)
2 {
3     py::class_<MatrixS>(m, "MatrixS")
4         .def(py::init<int, int>())
5         .def_readwrite("m", &MatrixS::m)
6         .def_readonly("n", &MatrixS::n)
7         .def_readonly("mat", &MatrixS::mat)
8         .def("randomize", &MatrixS::randomize)
9         .def("getitem", &MatrixS::getitem)
10        // .def("__getitem__", &MatrixS::getitem), py::is_operator())
11        .def("__repr__", &MatrixS::to_string)
12        .def(py::self + py::self)
13        .def(double() * py::self)
14        .def(py::self * double())
15    ;
16 }
```



# Practical Applications: quick prototype, pytest, input automation, visualization ...

## [McMillan He simulation](#) GitHub repository “McMillanHe”

1. C++ core class McMillanHe performs variational Monte Carlo random walk
2. Fortran analysis module grsk.f90 performs analysis on random walk samples for  $g(r)$  and  $S(k)$
3. Python workflow script mmh.py coordinate simulation, analysis, and visualization.



```
[yyang173@localhost McMillanHe]$ pytest -v .
===== test session starts =====
platform linux2 -- Python 2.7.15, pytest-3.7.4, py-1.6.0, pluggy-0.7.1 -
- /usr/bin/python2
cachedir: .pytest_cache
rootdir: /home/yyang173/Desktop/pyglue/mcmillan/nonboost/McMillanHe, ini
file:
collected 6 items

tests/test_mmh.py::test_init PASSED [ 16%]
tests/test_mmh.py::test_get_lbox PASSED [ 33%]
tests/test_mmh.py::test_set_lbox PASSED [ 50%]
tests/test_mmh.py::test_get_al PASSED [ 66%]
tests/test_mmh.py::test_set_al PASSED [ 83%]
tests/test_mmh.py::test_wfval PASSED [100%]

===== 6 passed in 0.11 seconds =====
```

Now we can use [Sphinx](#) instead of doxygen!



# References and Acknowledgement

## f2py

- [1] [f2py user guide](https://docs.scipy.org/doc/numpy/f2py) <https://docs.scipy.org/doc/numpy/f2py>
- [2] [scipy endorsement](#)
- [3] [illustrative example](#)
- [4] [our very own Katy](#)

## pybind

- [1] [pybind11 readthedocs](#)
- [2] [boost.python projects](#)
- [3] [pybind vs. swig](#)

## tutorial repositories

- [1] [basic steps](#)
- [2] [McMillan He simulation](#)

Many thanks to [Ryan Levy](#) for guiding me away from my boost.python debacle in favor of pybind11, which plays nicely with [Eigen](#).



## Conclusions:

- ❖ Thanks to packages like **f2py** and **pybind**, it is exceedingly easy to call C and Fortran from Python.
- ❖ However, SIMD and linear algebra can (should?) be left to numpy.
- ❖ Nested for loops can speed up  $O(10-100)$  times in C or Fortran.

**TL;DR** Flex your F/C muscles when you have to, but leave most complexity to your *Smart Pythonic “Brain”*.

